

# Eight agents. One closed loop. Evaluate, critique, improve — until the evidence says stop.

Describe a clinical case scenario. A live Planner → Executor → Evaluator → Critic → Improver cycle (with Memory and Safety oversight) drafts and refines a differential-diagnosis writeup for clinicians, calling the Groq API fresh every round, until a real stop condition fires.

⚠ EDUCATIONAL TOOL ONLY. Output is AI-generated differential-diagnosis support content intended for clinician review, not a diagnosis and not a substitute for clinical judgment, physical examination, or established guidelines. Do not use for



AUTONOMOUS AGENT STUDIO  
clinical differential · multi-agent pipeline

complete

CLINICAL CASE SCENARIO

58-year-old man with type 2 diabetes presents with 2 days of progressive left flank pain, fever (39.1°C), dysuria, and new confusion noted by family. Exam: hypotensive (92/58), tachycardic (118), CVA tenderness on left. Labs pending. No prior GU history.

TARGET SCORE THRESHOLD (STOP WHEN CROSSED)

8 / 10 — standard

HARD ITERATION CAP (SAFETY FALLBACK ONLY)

8 rounds max

GROQ API KEY

gsk\_...


Run pipeline

Reset

Calls `api.groq.com/openai/v1/chat/completions` — needs your own Groq API key, kept only in this browser tab (never sent anywhere else).

# Live orchestration dashboard

ROUND-BASED, STOP-CONDITION GOVERNED



Complete — stopped on: threshold

PIPELINE COMPLETE

Score 8.3 crossed the target threshold of 8.

|                 |                     |
|-----------------|---------------------|
| 1<br>ROUNDS RUN | 8.3<br>LATEST SCORE |
| 1<br>RETRIES    | 7<br>API CALLS      |

Activity log

---

11:16:58 pm **Critic** Critique produced.

---

11:16:58 pm **Safety Monitor** Running safety screen on current draft.

---

11:17:01 pm **Safety Monitor** Clear.

---

11:17:01 pm **System** Stop check after round 1: STOP (threshold).

---

11:17:01 pm **Final Reviewer** Composing final review based on stop reason and score history.

---

11:17:01 pm **System** Final Reviewer call failed (attempt 1/3): HTTP 429 – Rate limit reached for model `llama-3.1-8b-instant` in organization `org\_01k6x0f1z2ehzbcq9tvfrth07s` service tier `on\_demand` on tokens per minute (TPM): Limit 6000, Used 5217, Requested 1267. Please try again in 4.84s. Need more tokens? Upgrade to Dev Tier today at

---

## Iteration history

EVALUATION → CRITIQUE → IMPROVEMENT, PER ROUND

ROUND 1

8.3 / 10

---

### EVALUATOR — SCORE & REASONING

This draft differential diagnosis provides a clear leading diagnosis of sepsis, supported by the patient's hypotension, tachycardia, fever, and potential renal involvement. However, it falls short in clinical accuracy/evidence-based rigor, as it lacks explicit consideration of evidence-based guidelines for sepsis recognition and management. The draft does an excellent job of completeness of differential, including a range of reasonable considerations such as pyelonephritis, renal colic, and metabolic complications related to type 2 diabetes. The clarity and structure of clinical reasoning are well-organized and easy to follow, making it usable by a clinician under time pressure. However, the draft could improve in the area of safety by more explicitly addressing uncertainty and alternative explanations for the patient's presentation, particularly given the absence of previous GU symptoms and classic DKA symptoms.

---

### CRITIC — GAPS IDENTIFIED

Here are 7 concrete, specific, and actionable gaps for the Improver to address in the next round:

- \*\*Missing a specific criteria for sepsis recognition\*\*:** The draft mentions sepsis as the leading diagnosis but doesn't explicitly mention the sepsis 3 criteria (suspected infection, systemic inflammatory response syndrome (SIRS), and organ dysfunction) used for sepsis recognition. Ensure that the Improver includes these criteria in the draft.
- \*\*Lack of discussion on the role of type 2 diabetes\*\*:** The draft mentions the patient's type 2 diabetes under the differential for metabolic complications but doesn't consider how diabetes may affect the presentation and management of sepsis. The Improver should discuss this in more detail, considering the patient's underlying condition that may influence the severity of sepsis.
- \*\*Inconsistent inclusion of differential options\*\*:** While the draft includes a range of differential options,

the inclusion of "Metabolic complications related to type 2 diabetes" is not consistent with the absence of prior GU symptoms and classic DKA symptoms. The Improver should reconcile or exclude these differential options to ensure clinical accuracy.

4. **\*\*Lack of explicit consideration of alternative explanations\*\***: The draft mentions opposing evidence for sepsis but doesn't explicitly consider alternative explanations for the patient's presentation, such as other sources of infection or underlying conditions that may mimic sepsis (e.g., vasoplegic shock). The Improver should address this by explicitly considering alternative explanations.

5. **\*\*Insufficient consideration of safety and uncertainty\*\***: While the draft mentions red flags and pitfalls to watch for, it doesn't explicitly address the uncertainty associated with the patient's presentation. The Improver should incorporate a more nuanced discussion of safety and uncertainty, including the potential consequences of over- or under-treating sepsis.

6. **\*\*Missing recommended laboratory tests\*\***: Although the draft recommends broad-spectrum antibiotics and ordering urinalysis and imaging studies, it doesn't mention the necessity of ordering blood cultures and lactate levels to diagnose and manage sepsis. The Improver should include these laboratory tests in the recommended next steps.

7. **\*\*Inadequate consideration of organ-specific dysfunction\*\***: While the draft mentions the patient's confusion and potential renal involvement, it doesn't consider how sepsis may affect other organs, such as the liver, lungs, or adrenal glands. The Improver should discuss the potential for multi-organ dysfunction and how to diagnose and manage this complication in the context of sepsis.

---

#### MEMORY MANAGER — RUNNING NOTE

*Pending — runs after this round's Improver call, so it can compare the before/after draft rather than assume fixes landed.*

---

#### SAFETY MONITOR

✓ **CLEAR** — This draft appropriately acknowledges uncertainty by providing differential diagnoses and weighing supporting and opposing evidence for each option. It effectively flags genuine emergencies, including sepsis and acute abdominal conditions, and recommends immediate attention and treatment. Additionally, it suggests thorough evaluations and testing to identify the underlying cause of the patient's symptoms, demonstrating a cautious and thorough approach to diagnosis.

► View draft produced this round ▼

---

## Final summary

POPULATED WHEN A STOP CONDITION FIRES

STOP REASON: THRESHOLD

Score 8.3 crossed the target threshold of 8.

**\*\*Final Clinical Review\*\***

After thorough evaluation, the attached differential diagnosis for the patient's hypotensive and tachycardic state, fever, and potential renal involvement is deemed ready for clinician use. The document provides a comprehensive analysis of the patient's presentation, leading diagnosis, and potential differential diagnoses, including sepsis, pyelonephritis, renal colic, and metabolic complications related to type 2 diabetes.

While the score history meets the target threshold, clinician judgment remains crucial in interpreting the patient's symptoms and determining the best course of treatment. Key takeaways include the importance of immediate attention to sepsis symptoms, the need to rule out other emergency conditions, and the importance of comprehensive evaluation through urinalysis, blood cultures, imaging studies, and assessment of metabolic status.

Clinicians should be aware of potential pitfalls, including delayed assessment of CNS symptoms and failure to recognize sepsis. This document is decision-support content intended to aid clinician judgment, not to replace it. As such, it is essential to carefully evaluate the patient's presentation and adjust treatment accordingly.

**\*\*Conclusion\*\***

The differential diagnosis is confirmed as ready for clinician use. Clinicians should remain vigilant, taking into account any residual caveats and exercising caution when interpreting the patient's symptoms. This conclusion is based on the final draft and score history, meeting the target threshold of 8.

FINAL DIFFERENTIAL DRAFT

**\*\*Draft Differential Diagnosis\*\***

**\*\*1. Leading Diagnosis:\*\*** Sepsis

Given the patient's hypotensive and tachycardic state (92/58, 118 bpm), fever (39.1°C), and potential renal involvement (CVA tenderness on left), sepsis is a high-priority concern that requires immediate attention. The patient's presentation suggests a systemic infection that may have originated in the genitourinary (GU) system or elsewhere, prompting a comprehensive evaluation.

**\*\*2. Differential List:\*\***

1. **\*\*Pyelonephritis\*\***:

- Supporting evidence: Fever, dysuria, and CVA tenderness on left, suggesting a GU infection.

|   |           |      |   |
|---|-----------|------|---|
| 1 | 8.3 → 8.3 | +0.0 | 7 |
|---|-----------|------|---|

ROUNDS RUN

SCORE ARC

NET IMPROVEMENT

TOTAL API CALLS

## Agent performance summary

| AGENT          | INVOCATIONS THIS RUN |
|----------------|----------------------|
| Planner        | 1                    |
| Executor       | 1                    |
| Evaluator      | 1                    |
| Critic         | 1                    |
| Improver       | 0                    |
| Memory Manager | 0                    |
| Safety Monitor | 1                    |
| Final Reviewer | 1                    |

## Execution stats

7

API CALLS

1

RETRIES

1

ROUNDS

0

SAFETY FLAGS  
RAISED

## Architecture overview

HOW THE STOP-CHECK GOVERNS INFORMATION FLOW

Each round is a real, sequential API round-trip chain: **Evaluator** scores the current draft against a fixed rubric, **Critic** receives that score plus its reasoning and finds concrete gaps, and **Improver** receives the draft, evaluation, and critique together to produce the next draft. After Improver returns, **Memory Manager** distills what changed into a running memory note that is threaded into the next round's Improver call, and **Safety Monitor** screens the new draft for unsafe or overconfident clinical claims before the loop is allowed to continue. The stop-check runs after every round in strict order — plateau, then threshold, then hard cap — and its result determines whether another Evaluator call fires or the loop exits to **Final Reviewer**.

#### PLANNER

Reads the case once and sets structure/scope for the Executor. Runs only in round 1.

#### EXECUTOR

Writes the first full differential draft from the plan. Runs only in round 1; later rounds hand this job to Improver.

#### EVALUATOR

Scores the current draft 0-10 against a fixed 4-part clinical rubric, every round, with written reasoning.

#### CRITIC

Turns the Evaluator's reasoning into concrete, actionable gaps the Improver must address next.

#### IMPROVER

Rewrites the draft using the score, critique, and running memory note. Its output becomes next round's draft.

#### MEMORY MANAGER

Distills what changed this round into a short note carried forward so fixes aren't undone or repeated.

#### SAFETY MONITOR

Screens every new draft for unsafe overconfidence or missed red flags before the loop is allowed to continue.

#### FINAL REVIEWER

Runs once, after a stop condition fires, to close out the case with residual caveats for the clinician.

## Extension ideas

WHERE THIS PIPELINE COULD GO NEXT

### Guideline retrieval grounding

Give Executor a web-search or MCP tool call against UpToDate-style or PubMed sources so differentials are grounded in retrieved evidence rather than parametric knowledge alone, with citations carried through to Final Reviewer.

### Multi-case memory bank

Persist Memory Manager notes across sessions (e.g. via the Artifacts storage API) so recurring diagnostic blind spots the Critic keeps flagging get surfaced automatically on future cases.

### Human-in-the-loop gate

Insert an approval step before Final Reviewer where a clinician can accept, reject, or annotate the draft — turning the autonomous loop into a supervised one for higher-stakes deployments.

### Parallel specialist critics

Run Critic as three parallel calls (infectious disease lens, cardiovascular lens, red-flag/can't-miss lens) and merge their outputs before a single Improver pass, trading latency for broader coverage per round.

# Further prompts for building more advanced autonomous systems

|                                                                                                                      |      |
|----------------------------------------------------------------------------------------------------------------------|------|
| Add a human-approval gate before Final Reviewer that lets a clinician edit the memory note directly.                 | Copy |
| Extend the Evaluator rubric to a 6-criterion weighted score and show per-criterion sub-scores in the round card.     | Copy |
| Add a parallel 'red-team' Critic that argues for the single most dangerous missed diagnosis each round.              | Copy |
| Wire Executor to a real web-search tool for guideline grounding and require Final Reviewer to check citations exist. | Copy |

Autonomous Agent Studio – orchestration runs live against the Groq API. Educational demonstration;  
not a clinical decision-support device.