

Sixty days of building became **one thing shipped.**

The skill progression behind MedGuide AI — an agentic clinical research assistant, deployed live at the end of a 10-day capstone sprint.



WEEKS 1-2

AI Foundations & Effective Prompting

Learning how these models actually reason — and how to ask for what you want clearly enough that they can give it to you.



WEEKS 2-3

Research Workflows & Automation

Turning one-off prompts into repeatable processes — the first taste of building tools instead of just having conversations.



WEEKS 3-5

Product Thinking & Increasingly Sophisticated Projects

Moving from "can I build this?" to "should I build this?" — scope, users, and tradeoffs, not just capability.

○ WEEKS 5–7

Multimodal AI & Knowledge Systems

Working with more than text — and learning to ground AI output in real, retrievable documents instead of memory alone.

○ WEEKS 7–8

Autonomous Agent Architectures & Multi-Agent Collaboration

State machines, tool-calling, conditional routing — systems that decide what to do next, not just what to say next.

10-DAY CAPSTONE · THE PROOF

MedGuide AI — every skill above, shipped as one working product.

An agentic RAG assistant that answers clinical questions with cited sources and validates lab values against reference ranges — built through a full SDLC, from requirements to a live, production-hardened v1.0.0.



LangGraph agent orchestration

RAG · Chroma · Embeddings

Groq LLM

Production hardening

Live deployment

10

DAYS, REQUIREMENTS TO
PRODUCTION

302

DOCUMENT CHUNKS INDEXED

11

production issues found & fixed

V1.0.0

shipped & publicly live

MedGuide AI — built with Claude, AB Talks 60-Day Challenge capstone

github.com/Nidhisingh274/medguide-ai